

The Ethical Implications and Negative Consequences of Digital Assistants

Harlene Gill

Norquest College

CMPT 2110 H01

Kevin Laxamana

Position Paper

June 24, 2025

Abstract

Digital assistants (DAs) have woven themselves into the fabric of modern existence. These devices, driven by powerful language models, offer the illusion of companionship and intelligence while simplifying daily tasks. However, beneath this surface lies an intricate ethical web. From the erosion of privacy and autonomy to the simulation of emotional states and the silent reconfiguration of human values, DAs operate with little resistance and less scrutiny. This paper argues that while DAs offer undeniable convenience, they simultaneously raise significant ethical concerns that compromise our humanity, cognition, and trust. Using interdisciplinary evidence from philosophical, technological, and psychological studies, the following discussion will explore the many reasons we must reexamine our relationship with digital assistants before convenience replaces conscience.

The Ethical Implications and Negative Consequences of Digital Assistants

We now live in a world where speaking to a machine feels as ordinary as ordering coffee. DAs like Alexa, and ChatGPT respond to our queries, schedule our appointments, and even offer well-timed jokes. Once a tale of science fiction stories, these tools now exist everywhere from our homes to daycares, and professional settings. Their capabilities are impressive, but their impact extends far beyond functionality. In simulating human interaction, DAs subtly reshape the way we connect with others, absorb information, and interpret emotions. Their design, often framed as intuitive or empathetic, masks a reality where ethical implications are neither intuitive nor benign.

The stakeholders in this discussion are manifold: everyday users who unknowingly share data; vulnerable individuals like children or seniors who may place undue trust in these devices; corporations who profit off behavioral data; and policymakers who scramble to keep up. As society leans more heavily on AI-driven systems, it becomes imperative to understand the consequences of these technologies, not just from a technical standpoint, but an ethical one. The ethical dilemmas surrounding DAs are not theoretical, but realistic and urgent, requiring thorough consideration.

To help frame these issues we consider two dominant ethical frameworks: deontology and utilitarianism. Deontology emphasizes the importance of duties and moral rules. From this view, technologies must respect core principles like honesty, transparency, and consent. DAs often fall short here, simulating emotional responses without sincerity and collecting data without clear, informed user approval. On the other hand, utilitarianism evaluates actions by their outcomes, focusing on whether a technology maximizes well-being for the greatest number. While DAs may offer benefits like convenience and assistance for not only the vulnerable, but

rather the general public, they also pose widespread harms, especially through privacy loss, misinformation, and emotional detachment. These contrasting lenses reveal the moral complexity of relying on such devices in everyday life.

Digital assistants, while convenient, raise many ethical concerns that undermine autonomy, privacy, humanness, and epistemic integrity. Their integration into daily life poses threats that outweigh their benefits, demanding ethical scrutiny and diligent safeguards. Without protective policies, these technologies risk normalizing a distorted standard of interaction, knowledge, and emotional exchange. To understand the full scope of these risks, we must examine how DAs effect users on both an individual and societal level, beginning with the emotional and relational shifts they provoke through simulated sincerity.

A defining feature of DAs is their ability to mimic human emotional expression. These machines can express gratitude, offer apologies, and even convey romance in ways that can be alarming. On the surface, this might like a fun design choice meant to enhance user experience. However, this design crosses into ethically ambiguous territory when examined through a deontological lens. These devices present themselves as emotionally responsive agents while lacking any real capacity for empathy or care. Porra et al. (2019) emphasize this gap, stating, “a digital assistant cannot meaningfully produce expressive speech acts, and thus, be sincere” (p. 2) because they lack genuine psychological states. The deception lies not in the words themselves, but in the moral implication behind them. When a machine expresses statements such as “I understand how you feel”, they clearly cannot; yet it is this precise anthropomorphic empathy we crave and thus assign to objects we interact with. This misalignment between the original intent and the reality erodes the integrity of emotional exchanges and risks redefining our

understanding of sincerity. From a deontological standpoint, this is ethically immoral, as it fails to treat users as rational agents deserving of truth and respect.

What is even more concerning is how DAs threaten our very understanding of humanness. Porra et al. (2019) delve into this, arguing that our humanness is not a fixed trait, but something that emerges from shared emotional and social experiences over time. They explain that “humans bond systematically, automatically and subconsciously with other humans to form what she calls [Porra] calls human colonies...and that our humanness is the result of this kind of systematic bonding over very long time periods (Porra et al., 2019, p. 3). This framing is crucial. DAs do not just imitate interaction, they simulate the emotional cues that form the foundation of human identity. Without the ability to feel or reciprocate, DAs reduce emotional bonding to feature performance. Worse, Porra et al. (2019) warns that “we don’t really know what we are replicating because of the unanswered questions surrounding genuine humanness” (p.4). The result is an interaction that feels human without being human, and that subtle difference carries ethical weight. If we continue to engage with these systems as if they are part of our emotional ecosystem, we risk transforming ourselves in the process. What begins as harmless mimicry could shift how we experience connection or feel empathy. From both deontological and utilitarian perspectives, that potential for slow erosion of human depth and symbolic meaning must be taken seriously.

Beyond emotional integrity, DAs also impact the way we come to know and understand the world. They answer questions with confidence and immediacy, often without disclosing limitations or sources of their knowledge. This dynamic creates a subtle but strong epistemic source. Freiman (2023) describes the emerging category of “technology-based beliefs”, where users acquire beliefs not from human-testimony but from non-human agents; a shift that lacks

traditional avenues for scrutiny or correction. While people defending utilitarianism may argue that accuracy is key and enough for the long-term, this approach ignores the risks of systemic misinformation and user dependency. From a deontological view, allowing machines to present unverified or biased information as truth violates our duty to ensure absolute the absolute truth. It reduces the user's ability to question, challenge, and discern both moral agency and intellectual independence.

Freiman's work raises an important question: "does anyone hold any responsibility for these devices' truth inputs and outputs?" (Freiman, 2023, p.3). Unlike traditional human testimony, there is no accountable agent behind a DAs' output. No face, nor context, just information, presented confidently. This lack of responsibility means users are left to assume reliability. As Freiman (2023) notes, "when an instrument repeatedly produces a truth output, the subject is inclined to incorporate its reliability as an assumption" (p. 6). The more we hear the right answer, the less we question where it came from. This slowly builds an unearned trust. What is more, Frieman (2023) explains that existing epistemic models often focus on the reliability of the tool, not the actual content of what is being delivered or its origin (p. 6). So not only do users struggle to evaluate what is being said, they also lack the tools to question it properly. This creates a real risk; beliefs are formed without scrutiny, and knowledge becomes detached from justification. From both a utilitarian and deontological view, the damage is clear, users are no longer fully autonomous in how they form beliefs. They are being nudged toward passive acceptance, a dangerous ethical dilemma.

Privacy too, is a central ethical battleground. DAs are designed to be ready, constantly listening for a "power on" word but often hearing more than intended. Simon et al. (2022) present their argument through a scenario-based approach. Here they create a "near-future

vignette-a digital assistant that has capabilities to detect early signs of cognitive impairment (Simon et al, 2022, p.1). This device, like most DAs currently in the market (ex. Alexa), was most “likely to have been offered a click-through user agreement” (Simon et al., 2022, p.2) that obscures how her data is collected and used. While such tools may promise convenience and even health benefits, their mechanisms bypass traditional consent models, reducing deeply personal data into a commercial asset. This setup violates both ethical frameworks: from a deontological perspective, it treats individuals as a means to data collection rather than ends in themselves; from a utilitarianism view, the combined harm of undermined trust and constant surveillance outweighs short-term convenience.

The values embedded in DAs extend beyond privacy and knowledge, they also shape our identities and moral judgements. Kudina (2021) explains how DAs recontextualize meaning for users, noting that “such design features...help to shape them [users] as specific subjects” (p. 2) through daily interaction. These technologies are not neutral; they guide users’ perspectives, decisions, and emotional responses. For instance, the persistent use of female-sounding assistants subtly reinforces gendered assumptions about catering to CIS gendered males and the emotional labor required of relationships with patriarchal views. This social coding, embedded in technology, is neither accidental nor harmless. Through the lens of deontology, it imposes a moral narrative without user awareness or consent. Utilitarianism, too, reveals cracks: any benefit gained from ease of interaction comes at the cost of perpetuating outdated and harmful social norms.

Even the promise of connection, a key selling point for many DAs, warrants skepticism. Candiotto (2022) discusses the modern paradox of digital connection, stating that extended loneliness does not stem from an actual lack of connections, but of longing for meaningful

relationships while being connected to many people online. DAs may mimic companionship, but they cannot reciprocate it. During moments of isolation, such as the COVID-19 pandemic, many turned to their devices for a sense of connectivity. Candiotto (2022) points out this connectivity may have its moments, but overall creates a hollow facsimile of a relationship. This substitution risks redefining what it means to feel understood or cared for, blurring the emotional expectations we place on both humans and machines. In this context, DAs not only fail to alleviate loneliness, but they may also deepen it by masking its roots.

Critics may argue that DAs promote independence, especially for the elderly or those with disabilities. They enable users to manage their routines, communicate with caregivers, and maintain a level of self-sufficiency that might otherwise be impossible. From a utilitarian view, these benefits are significant and worth preserving. Similarly, proponents of “ethics by design” (Simon et al., 2022, p.1) suggest that better regulatory frameworks and user controls can mitigate most ethical concerns. If users are given more transparent consent options, if devices are built to reflect ethical values from the ground up, then perhaps the harms can be reduced.

These arguments rest on shaky ground. The reality is that the current digital ecosystem is driven by profit, not ethics. Ethical design, while significant, remains the exception rather than the norm. Transparency is undermined by legal complexity and vague documentation. Users may technically have options, but they often lack the understanding or time to exercise them meaningfully. Additionally, the anthropomorphism that supposedly improves accessibility also blurs critical distinctions between tools and companions. By encouraging users to treat machines as if they were people, DAs distort social expectations and emotional norms. Independence is valuable, but not when it comes at the cost of truth, privacy, and authentic connection.

Looking ahead, the ethical concerns surrounding DAs are likely to intensify. As these technologies become more sophisticated, integrating emotional AI, behavioural prediction, and even therapeutic functions, the line between human and machine interaction will blur further. Without rigorous ethical frameworks and robust regulatory oversight, DAs risk becoming not just tools of convenience, but instruments of manipulation. Their influence will extend into education, healthcare, governance, and more. Both deontological and utilitarian theories highlight the need for proactive engagement: the former demands respect for autonomy and sincerity; the latter insists on weighing long-term consequences over short-term comfort. The future of these technologies depends on our willingness to ask difficult questions now.

Conclusion

DAs are no longer simple tools; they are active participants in reshaping how we understand the world. What may begin as harmless convenience quickly turns into dependency on systems that simulate care, obscure truth, and redefine connection. As this paper has argued, the ethical risks embedded in their design are far from neutral. From the manipulation of emotional expression to the distortion of knowledge and erosion of privacy, DAs challenge foundational principles of both deontological and utilitarian ethics. They fail to respect our autonomy and sincerity, while also producing long-term harms that outweigh their short-term benefits. If left unchecked, their influence will continue to blur the lines between authentic human experience and artificial mimicry. It is not enough to marvel at their capabilities; we must understand their consequences. As digital assistants become increasingly woven into daily life, society must take responsibility in setting boundaries, demanding transparency, and reasserting what it means to be human in a world that is becoming saturated with digital technologies.

References

- Simon, D.A., Evans, J. B., Shachar, C., Cohen, G. (2022). Should Alexa diagnose Alzheimer's? Legal and ethical issues with at-home consumer devices. *Cell Reports Medicine* 3, 100692. <https://doi.org/10.1016/j.xcrm.2022.100692>
- Kudina, O. (2021). "Alexa, who am I?": Voice Assistants and Hermeneutic Lemniscate as the Technologically Mediated Sense-Making. *Human Studies* 44:233-253. <https://doi.org/10.1007/s10746-021-09572-9>
- Freiman, O. (2023). Analysis of Beliefs Acquired from a Conversational AI: Instruments-based Beliefs, Testimony-based Beliefs, and Technology-based Beliefs. *Episteme*, 21 (1031-1047). doi:10.1017/epi.2023.12
- Porra, J., Lacity, M., Parks, M.S. (2019). "Can Computer Based Human-Likeness Endanger Humanness?" – A Philosophical and Ethical Perspective on Digital Assistants Expressing Feelings They Can't Have. *Information Systems Frontiers*, 22:533-547. <https://doi.org/10.1007/s10796-019-09969-z>
- Candiotto, L (2022). Extended loneliness; When hyperconnectivity makes us feel alone. *Ethics and Information Technology*, 24:47. <https://doi.org/10.1007/s10676-022-09669-4>